



Advanced Analytics with R

Text Analytics

Text Analytics

- *Text analytics*, as known as *text mining*, is the process of obtaining patterns and trends from text.
- Typical text analytics tasks include
 - text categorization
 - text clustering
 - concept/entity extraction
 - production of granular taxonomies
 - *sentiment analysis*
 - document summarization, and
 - entity relation modeling

Applications of Text Analytics



Applications of Text Analytics



Applications of Text Analytics



Getting Twitter Data

- <https://sites.google.com/site/miningtwitter/basics/getting-data/by-twitter>
- <https://www.credera.com/blog/business-intelligence/twitter-analytics-using-r-part-1-extract-tweets/>

Getting Facebook Data

- <https://bigdataenthusiast.wordpress.com/2016/03/19/mining-facebook-data-using-r-facebook-api/>

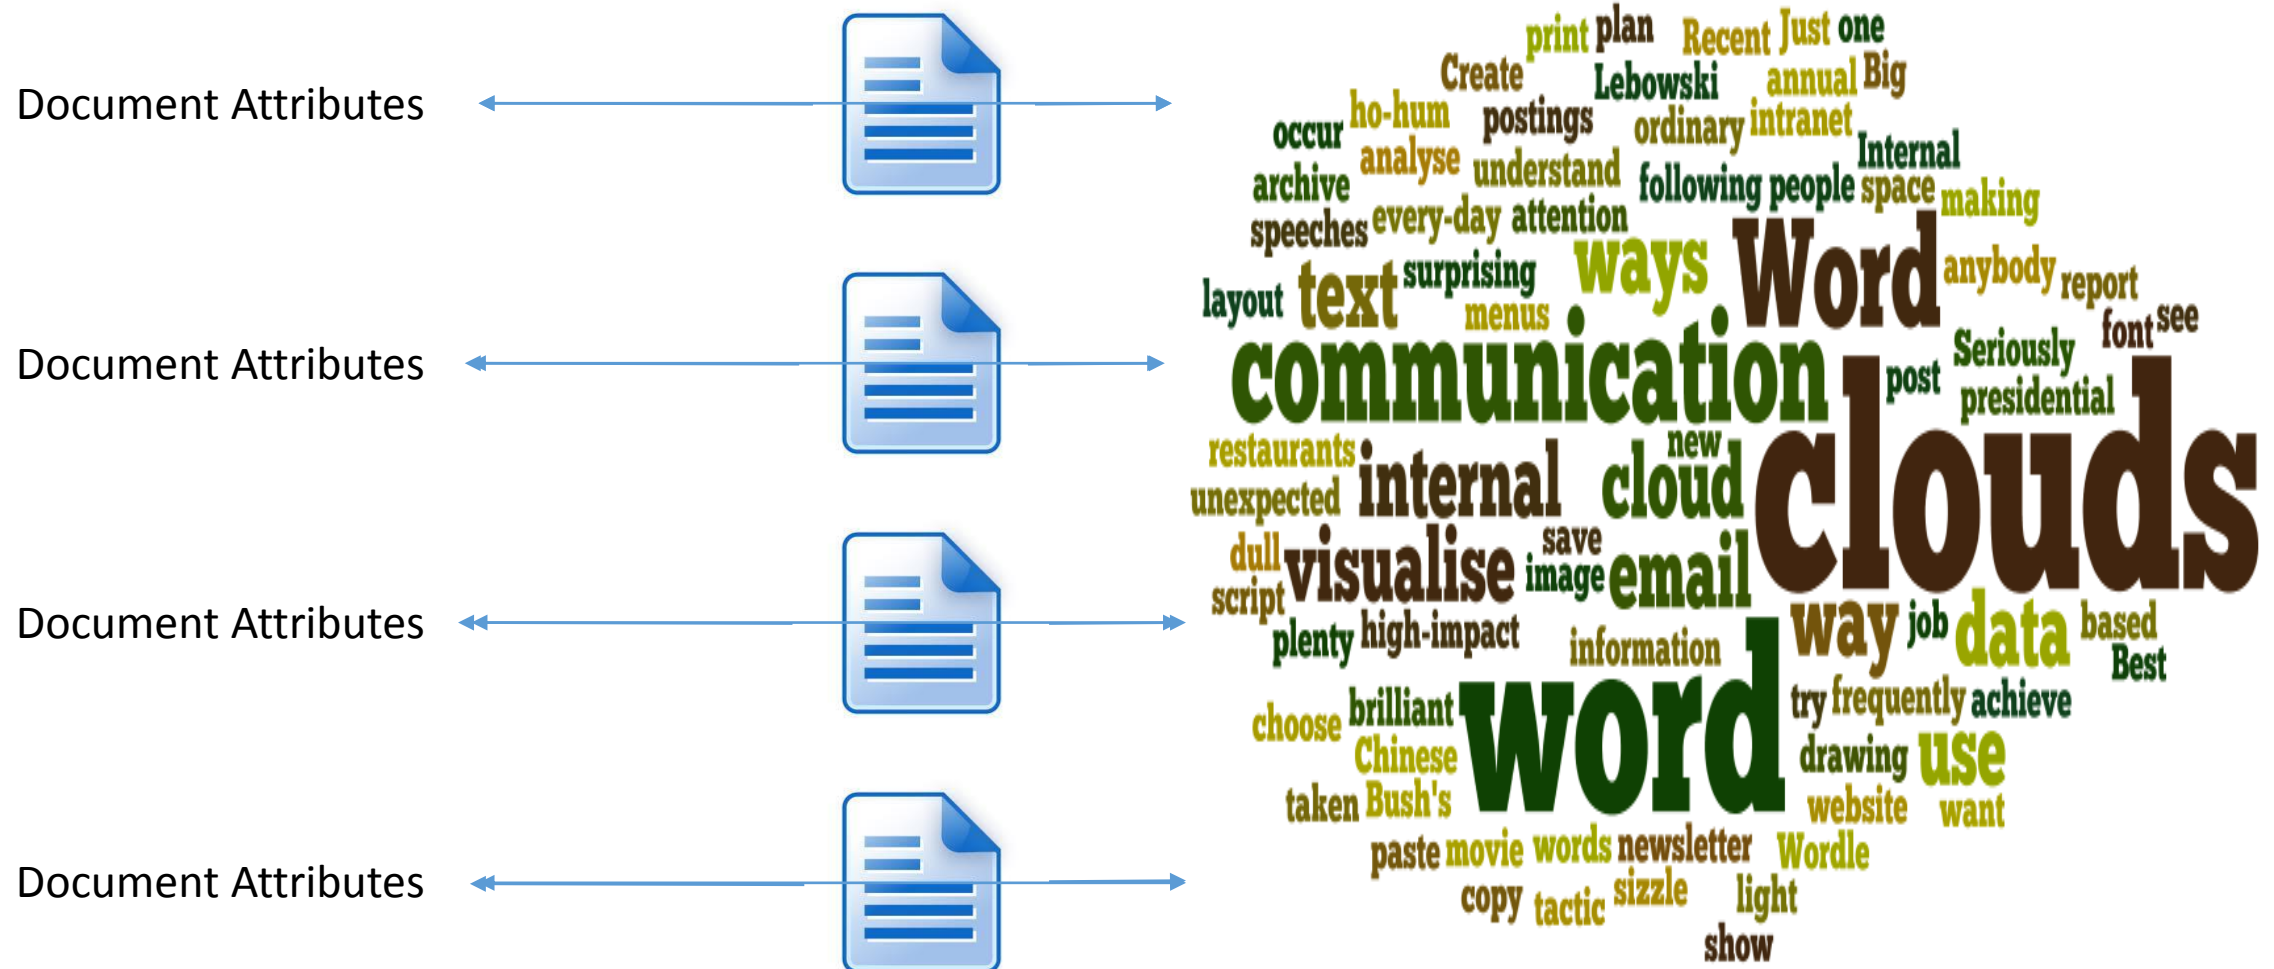
Getting Instagram Data

- <https://www.r-bloggers.com/analyze-instagram-with-r/>

Case – Movie Review

- Data: Reviews on 1000 movies from the Internet Movie Database (IMDb).
- Task: to study whether we can predict the movie rating or sentiment (thumbs up or down) from the reviews.

Basic Concept



Step 1. Data Preparation

Corpus

- Normally, texts are originally stored in various documents, each of which has its own attributes, such as date, owner, rating, etc.
- A collection of text documents are stored in a structure called **Corpus** for exploration.
- The default implementation is the so-called **VCorpus** (short for Volatile Corpus) which realizes a semantics as known from most R objects: corpora are R objects held fully in memory. We denote this as volatile since once the R object is destroyed, the whole corpus is gone.
- Another implementation is the **PCorpus** which implements a Permanent Corpus semantics, i.e., the documents are physically stored outside of R (e.g., in a database), corresponding R objects are basically only pointers to external structures, and changes to the underlying corpus are reflected to all R objects associated with it. Compared to the volatile corpus the corpus encapsulated by a permanent corpus object is not destroyed if the corresponding R object is released

Read in data

```
#Let's first load data stored in /reviews/train/unsup/, which are the unsupervised data  
directory.location <- paste(getwd(),"/reviews/train/unsup/",sep = "")
```

```
#Read in data and stored it as a Corpus  
movie.unsup.corpus <- Corpus(DirSource(directory.location, encoding = "UTF-8"),  
  readerControl = list(language = "en_US"))
```

Inspect Corpus

#check the number of documents stored in a corpus
`length(movie.unsup.corpus)`

#Have an overview of (part of) the corpus
`inspect(movie.unsup.corpus[1])`

#view content of a particular document in a corpus
`movie.unsup.corpus[[1]]$content`

#You can even use lapply function on corpus
`review.list <- lapply(movie.unsup.corpus, as.character)`

Data Transformation

- Steps are to be performed to transform documents to make the subsequent analysis more meaningful and more efficient:
 - Make the words more consistent
 - Good = good
 - Wants = want = wanted = wanting
 - Remove unnecessary words
 - White space
 - Punctuations
 - Numbers
 - Stop words

Further data transformation

- Identify top words in the corpus
- Inspect all the top words so that we can
 - Remove meaningless words
 - Correct key words

Word cloud

```
wordcloud(movie.unsup.corpus,max.words=100,random.order=FALSE,colors=brewer.pal(8, "Dark2"))
```



Sentiment Analysis

Load Positive and Negative Dictionary

```
bytecode.convert <- function(x) {iconv(enc2utf8(x), sub = "byte")}

# read in positive and negative word lists from Hu and Liu (2004)
positive.data.frame <- read.table(file = "Hu_Liu_positive_word_list.txt",
  header = FALSE, colClasses = c("character"), row.names = NULL,
  col.names = "positive.words")
positive.data.frame$positive.words <-
  bytecode.convert(positive.data.frame$positive.words)

negative.data.frame <- read.table(file = "Hu_Liu_negative_word_list.txt",
  header = FALSE, colClasses = c("character"), row.names = NULL,
  col.names = "negative.words")
negative.data.frame$negative.words <-
  bytecode.convert(negative.data.frame$negative.words)
```

Detecting Top Positive Words

```
Hu.Liu.positive.dictionary <-  
  c(as.character(positive.data.frame$positive.words))  
reviews.tdm.Hu.Liu.positive <- TermDocumentMatrix(movie.unsup.corpus,  
  list(dictionary = Hu.Liu.positive.dictionary))  
examine.tdm <- removeSparseTerms(reviews.tdm.Hu.Liu.positive, 0.95)  
top.words <- Terms(examine.tdm)  
print(top.words)  
Hu.Liu.frequent.positive <- findFreqTerms(reviews.tdm.Hu.Liu.positive, 25)  
# this provides a list positive words occurring at least 25 times  
# a review of this list suggests that all make sense (have content validity)  
# Then we accept Hu.Liu.frequent.positive fully as our positive dictionary  
test.positive.dictionary <- c(as.character(Hu.Liu.frequent.positive))
```


Detecting Top Negative Words

```
Hu.Liu.negative.dictionary <-  
  c(as.character(negative.data.frame$negative.words))  
reviews.tdm.Hu.Liu.negative <- TermDocumentMatrix(movie.unsup.corpus,  
  list(dictionary = Hu.Liu.negative.dictionary))  
examine.tdm <- removeSparseTerms(reviews.tdm.Hu.Liu.negative, 0.97)  
top.words <- Terms(examine.tdm)  
print(top.words)  
Hu.Liu.frequent.negative <- findFreqTerms(reviews.tdm.Hu.Liu.negative, 15)  
  
# this provides a short list negative words occurring at least 15 times  
# across the document collection...  
# By checking the list, we found that one of these words seems out of place  
# as they could be thought of as positive: "funny"  
test.negative <- setdiff(Hu.Liu.frequent.negative, c("funny"))  
# Now we exclude "funny" from the negative dictionary  
test.negative.dictionary <- c(as.character(test.negative))
```

From Text to Numbers

	Doc1	Doc2	...	Doc1000
Good	1	0		2
Amaze	0	1		0
⋮				
Favorite	1	0		1
⋮				
Fall	1	0		1
Dead	2	1		0
⋮				
Hell	0	0		3

TermDocumentMatrix



	Doc1	Doc2	...	Doc1000
#Total words	230	340		135
#Positive words	20	30		2
#Negative words	10	20		13
#Other words	200	290		120
Rating	10	5		2

Predictive Analysis

Simple classification

- Classify a movie into thumbs up (rating >5) or thumbs down (rating ≤ 5) by checking if there is more positive words or more negative words in the movie review.

Linear Regression

- Predict rating by linear regression on various numbers of words.
- Predicting thumbs up or down by using the linear regression result.

Further Study

- Applying other classification methods
- Weighted regression on various key words
- Cluster analysis